

# Eric (Ross) Fu

713-540-4528 | ericrossfu@yahoo.com | Houston – Open to Relocation (Los Angeles Miami Austin preferred) | [LinkedIn](#)

## SUMMARY

---

Data Scientist with 7.5 years of experience delivering high-impact data infrastructure and production-grade ML systems. Blends deep technical skill with business focus to deliver end-to-end value through data science, engineering, and MLOps.

## EDUCATION

---

***Southern Methodist University***  
Master of Science in Data Science

**Dallas, Texas**  
December 2020

***Texas A&M University***  
Bachelor of Science in Biology

**College Station, Texas**  
December 2015

## EXPERIENCE

---

Data Scientist 2021-Present  
Avance Biosciences Inc.

- Architected and deployed scalable, production-grade ML models on **AWS SageMaker**, often collaborating cross functionally to deliver impactful results on large-scale biological datasets and enterprise-grade data pipelines.
- Led the **end-to-end ML lifecycle**—from data engineering to model deployment—leveraging MLOps best practices (MLflow, versioning, governance, CI/CD) to ensure reproducibility, scalability, and traceability.
- Engineered large-scale **GenAI-powered Q&A system** using proprietary corpora, unlocking domain-specific insights and powering intelligent search, summarization, and recommendations across internal knowledge bases.
- Built XGBoost **churn model** with 85% recall, enabling retention team to reduce customer loss by 12%
- Used Deep Learning methods to perform **demand forecasting** for products with 92% accuracy, utilizing time-based features and improving supply planning.
- Built a real-time **Kafka → S3 → Glue pipeline** to stream data from 70+ sources into a partitioned, ML-ready data lake, enabling ETL and scalable model training.
- Designed and deployed a **real-time anomaly detection** solution leveraging Kafka streaming and Spark-powered AWS Glue to enable instant fault identification in lab equipment.
- Developed **computer vision**-enhanced ingestion systems that boosted data intake throughput by 10×, supporting automation of previously manual, bottlenecked processes.
- Doubled analytics throughput by scaling SQL infrastructure, unlocking company-wide self-serve reporting and eliminating data silos across teams.
- Delivered a **600× efficiency improvement** in core workflows by automating high-friction manual processes, reducing human error, and enabling real-time analytics through streamlined ML + data pipelines.
- Designed and analyzed A/B experiments to validate model performance in production, integrating statistical rigor into deployment pipelines.
- Technologies Used: Python, SQL, Machine Learning, Kafka, S3, Glue, Sagemaker, Lambda, Docker, AI, SAS

## SKILLS

---

### Languages & Tools:

Python, SQL, PySpark, SAS, R, Git, Linux\*

### Machine Learning & AI:

Deep Learning, Scikit-learn, boosted models, NLP, LLM, TensorFlow, PyTorch, HuggingFace, etc.

### Cloud & Data Infrastructure:

AWS (S3, SageMaker, Glue, Lambda, Kinesis), Kafka, Spark\*, Airflow\*, Databricks\*, FastAPI, Docker, Streamlit